

DIFFA: Large Language Diffusion Models Can Listen and Understand

Jiaming Zhou^{1*}, Hongjie Chen², Shiwan Zhao¹, Jian Kang², Jie Li², Enzhi Wang¹,
Yujie Guo¹, Haoqin Sun¹, Hui Wang¹, Aobo Kong¹, Yong Qin^{1†}, Xuelong Li^{2†}

¹ College of Computer Science, Nankai University,

² Institute of Artificial Intelligence (TeleAI), China Telecom, China,

Correspondence: zhoujiaming@mail.nankai.edu.cn, qinyong@nankai.edu.cn

Abstract

Recent advances in large language models (LLMs) have shown remarkable capabilities across textual and multimodal domains. In parallel, diffusion-based language models have emerged as a promising alternative to the autoregressive paradigm, offering improved controllability, bidirectional context modeling, and robust generation. However, their application to the audio modality remains underexplored. In this work, we introduce **DIFFA**, the first diffusion-based large audio-language model designed to perform spoken language understanding. DIFFA integrates a frozen diffusion language model with a lightweight dual-adaptor architecture that bridges speech understanding and natural language reasoning. We employ a two-stage training pipeline: first, aligning semantic representations via an ASR objective; then, learning instruction-following abilities through synthetic audio-caption pairs automatically generated by prompting LLMs. Despite being trained on only 960 hours of ASR and 127 hours of synthetic instruction data, DIFFA demonstrates competitive performance on major benchmarks, including MMSU, MMAU, and VoiceBench, outperforming several autoregressive open-source baselines. Our results reveal the potential of diffusion-based language models for efficient and scalable audio understanding, opening a new direction for speech-driven AI. Our code will be available at <https://github.com/NKU-HLT/DIFFA.git>.

1 Introduction

Large language models (LLMs) have catalyzed a paradigm shift in artificial intelligence, pushing the frontiers of natural language understanding, computer vision, and multimodal reasoning (Achiam et al., 2023). In the domain of speech and audio processing, large audio-language models (LALMs)




Figure 1: DIFFA vs. Qwen2-Audio-Instruct. The abbreviations correspond to MMSU benchmark’s capabilities: Perception-Semantics (Percep-Sem), Perception-Phonology (Percep-Phon), Perception-Paralinguistics (Percep-Paral), Reasoning-Semantics (Reason-Sem), Reasoning-Phonology (Reason-Phon), and Reasoning-Paralinguistics (Reason-Paral).

have similarly benefited from these advances in LLMs. By bridging continuous acoustic signals with discrete linguistic representations, LALMs enable end-to-end modeling of spoken interaction (Chu et al., 2023; Zhang et al., 2023). This capability not only advances fundamental research in speech understanding and generation, but also opens up practical opportunities for building more natural, robust, and versatile human-computer communication systems.

Existing LALMs typically follow two design paradigms. The first couples a speech encoder with an LLM, often through lightweight adapters that project continuous acoustic representations into the input space of the language model (e.g., Qwen2-Audio (Chu et al., 2024), Audio-Flamingo (Kong et al., 2024a)). The second discretizes audio into speech tokens via speech tokenizers and subsequently trains directly on these tokens under the LLM training paradigm (e.g., SpeechGPT (Zhang

* This work was done during an internship at TeleAI.

† Yong Qin and Xuelong Li are corresponding authors.

et al., 2023), Moshi (Défossez et al., 2024a)). Despite their strong results, both paradigms predominantly rely on autoregressive (AR) decoding, which suffers from well-known drawbacks such as exposure bias, slow generation, and limited flexibility for bidirectional or partially conditioned inference.

To address these limitations, diffusion-based language models (Austin et al., 2021; Shi et al., 2024) have emerged as a promising alternative. By framing generation as an iterative denoising process, diffusion models support non-autoregressive decoding, parallel prediction, and improved controllability (Shi et al., 2024). Recent advances such as LLaDA (Nie et al., 2025) demonstrate that diffusion LLMs can rival autoregressive counterparts like LLaMA-3 (Dubey et al., 2024), while exhibiting stronger robustness and training efficiency. Furthermore, LLaDA-V (You et al., 2025) extends this paradigm to vision-language tasks, confirming the competitiveness and generality of diffusion modeling in multimodal learning.

However, the audio modality remains notably underexplored in the context of diffusion-based language models. While such models have demonstrated promising results in text domains, their applicability to audio-language understanding has not been systematically investigated. The unique characteristics of audio—such as acoustic variability, complex temporal structures, and rich paralinguistic information—motivate an exploration into whether diffusion LLMs can be effectively extended to this domain with their flexible decoding mechanisms and bidirectional context modeling.

To bridge this gap, we explore the potential of adapting large diffusion-based language models for audio understanding. Specifically, we investigate whether such models can effectively process audio inputs and perform on par with, or surpass, strong autoregressive LALMs. Toward this goal, we introduce **DIFFA**, a **DIFF**usion-based large **A**udio-language framework. DIFFA adopts a modular and efficient design: a pretrained speech encoder (Whisper (Radford et al., 2023)), two lightweight adapters (semantic and acoustic), and a frozen diffusion-based LLM. To avoid catastrophic forgetting, we train only the adapters and keep both the language model and speech encoder frozen. The training procedure is divided into two stages: first, we align the semantic adapter under an ASR objective using LibriSpeech (Panayotov et al., 2015); then, we fine-tune both adapters on synthetic instruction data using the “What can you hear from

the audio?” prompting scheme inspired by the DESTA-2 (Lu et al., 2025). Figure 1 visually compares the performance of DIFFA and Qwen2-Audio-Instruct across multiple benchmarks. Notably, DIFFA attains competitive outcomes using merely 960 hours of ASR data and 127 hours of synthetic data, whereas Qwen2-Audio-Instruct depends on a far larger dataset of 510,000 hours.

Our contributions are summarized as follows:

- We propose the first diffusion-based LALM, DIFFA, enabling large-scale audio-text understanding without relying on autoregressive modeling.
- We introduce a dual-adapter training framework and a two-stage training strategy that aligns speech representations to a frozen diffusion LLM, supporting audio understanding and instruction following without explicit supervised fine-tuning data.
- Despite using only 960 hours of ASR data, 127 hours of synthetic instruction data and 72 A800 GPU hours, DIFFA achieves competitive performance across multiple benchmarks, including MMSU, MMAU, and VoiceBench.
- We will release the training pipeline, inference code, and data generation scripts to promote research on diffusion-based LALMs with minimal compute and data requirements.

2 Related Work

2.1 Large Audio-Language Models (LALMs)

Large audio-language models (LALMs) have recently emerged as powerful tools for spoken language understanding in open-ended tasks. Existing approaches generally fall into two paradigms. The first couples a speech encoder with a pretrained LLM, often via lightweight adapters to bridge acoustic and textual representations. Representative examples include Qwen-Audio (Chu et al., 2023) and Qwen2-Audio (Chu et al., 2024), which integrate Whisper through a unified interface; SALMONN (Tang et al., 2024a), which introduces a dual-encoder and window-level Q-Former to better capture long audio contexts; and Audio-Flamingo2 (Ghosh et al., 2025), which employs curriculum learning and dense alignment strategies. These systems achieve strong results on benchmarks such as MMAU and VoiceBench but remain tied to autoregressive decoding.

The second paradigm discretizes speech into tokens using quantizers or self-supervised encoders, and treats them as an additional input stream to the LLM (e.g., SpeechGPT (Zhang et al., 2023), Moshi (Défossez et al., 2024a)). While effective, this line of work also relies exclusively on autoregressive generation. To date, non-AR architectures such as diffusion models have not been explored for audio–language understanding.

2.2 Diffusion-Based Language and Multimodal Models

Diffusion-based models redefine generative modeling by reversing a corruption process that progressively masks or perturbs input data. Diffusion-LM (Austin et al., 2021; Shi et al., 2024; Sahoo et al., 2024) first applied this idea to discrete language modeling. LLaDA (Nie et al., 2025) subsequently scaled diffusion models to LLMs, showing strong performance on NLU and generation tasks. Notably, LLaDA employs a masked diffusion process with a principled likelihood objective and exhibits robustness in reasoning tasks. Its extension, LLaDA-V (You et al., 2025), introduces a visual pathway for multimodal generation, achieving results competitive with autoregressive vision-language models.

These results motivate our exploration of diffusion-based modeling for audio, a modality that naturally benefits from bidirectional and non-sequential reasoning. Our work represents the first application of this framework to LALMs.

2.3 Modality Alignment Without Explicit Supervised Fine-tuning

A growing body of work seeks to eliminate reliance on human-annotated SFT datasets. BLSP (Wang et al., 2023) and AudioChatLlama (Fathullah et al., 2024) introduce similar *behavior alignment*, where speech and text are treated as semantically equivalent inputs expected to elicit the same model output, but constrained in semantic alignment. DESTA (Zhang et al., 2023) and DESTA-2 (Lu et al., 2025) leverage synthetic instruction-tuning data, constructed by prompting LLMs with “What can you hear from the audio?” alongside speech metadata, to enable paralinguistic understanding. These strategies reduce the need for paired human supervision. However, DESTA-2 typically follow a cascaded design, relying on Whisper for transcription before instruction response. In contrast, our model integrates both semantic and acoustic infor-

mation via dual adapters and performs instruction-following in an end-to-end fashion using a frozen diffusion LLM.

3 Methods

In this section, we present the overall framework of our proposed **DIFFA**, including its formulation, data construction, training strategy, and inference procedure.

3.1 Preliminaries

LLaDA (Nie et al., 2025) is a non-autoregressive language modeling paradigm that introduces a discrete random masking process and learns a *mask predictor* to approximate its reverse. Unlike traditional autoregressive models, LLaDA allows for bidirectional dependency modeling and efficient likelihood-based training.

LLaDA defines a forward masking process to sample a corrupted sequence x_t , where each token is independently replaced with a special mask token M with probability $t \in (0, 1]$. The mask predictor $p_\theta(x_0|x_t)$ is parameterized by a standard Transformer decoder and trained to reconstruct masked tokens:

$$\mathcal{L}(\theta) \triangleq -\mathbb{E}_{t, x_0, x_t} \left[\frac{1}{t} \sum_{i=1}^L \mathbf{1}[x_t^i = M] \log p_\theta(x_0^i | x_t) \right], \quad (1)$$

where L denote the length of target sequence. This objective yields a tractable upper bound of the negative log-likelihood (Shi et al., 2024; Ou et al., 2025), enabling parallel token prediction.

Supervised fine-tuning (SFT) under LLaDA follows a similar approach. Given a prompt p_0 and response r_0 , the response tokens are masked independently to obtain r_t . The loss is computed as:

$$-\mathbb{E}_{t, p_0, r_0, r_t} \left[\frac{1}{t} \sum_{i=1}^{L'} \mathbf{1}[r_t^i = M] \log p_\theta(r_0^i | p_0, r_t) \right], \quad (2)$$

where L' is the response length.

During inference, LLaDA decodes iteratively from a fully masked sequence. At each denoising step, the model predicts masked tokens and re-applies masks to low-confidence positions, gradually refining predictions over T steps.

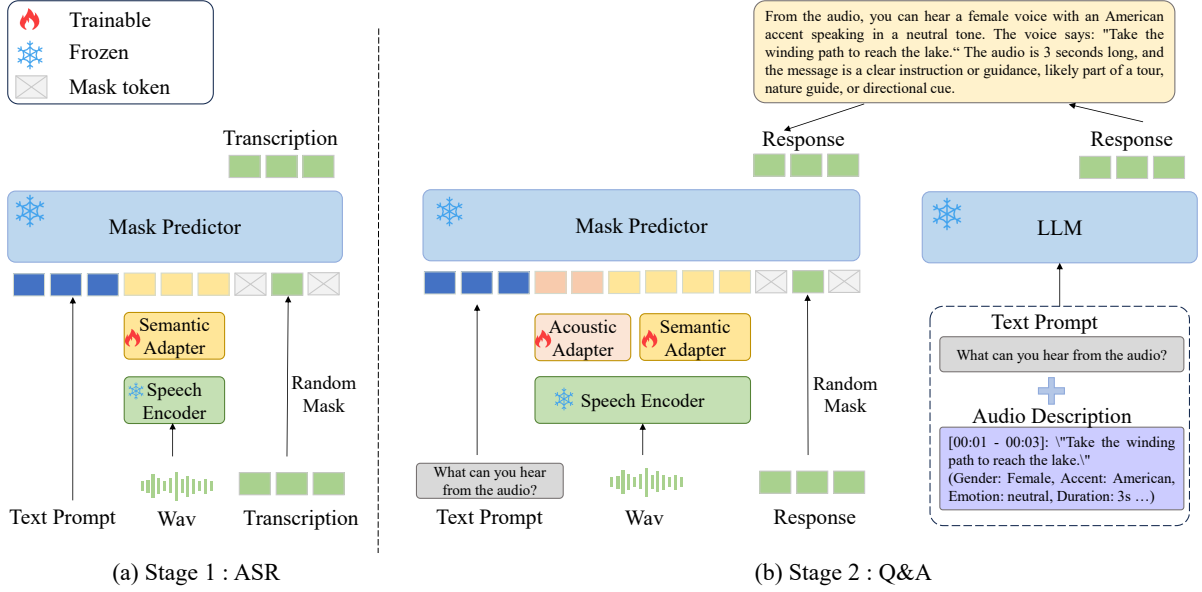


Figure 2: Training process of our DIFFA framework. Stage 1 performs semantic alignment via an ASR objective, aligning the speech encoder with the diffusion language model. Stage 2 enables modality alignment by prompting the model to describe what it hears from the audio, following an audio caption instruction paradigm.

3.2 Data Construction

Inspired by the DESTA series, we construct a dataset by prompting LLaDA-based or instruction-tuned language models (e.g., LLaMA3, Qwen3) with audio transcriptions: "[00:01 - 00:03]: *"Take the winding path to reach the lake."* (Gender: Female, Accent: American, Emotion: neutral, Duration: 3s ...) " and acoustic attributes using the prompt: "What can you hear from the audio?". The generated response serves as the supervision signal, paired with the corresponding audio. This enables extbmodality alignment without any explicit supervised fine-tuning data.

Furthermore, motivated by self-distillation techniques (Yang et al., 2024), we introduce a rewriting step to mitigate the domain shift arising from different model pre-training distributions. Specifically, we first employ Qwen3-8B to generate an initial set of captions. Subsequently, our LLaDA model rewrites these captions to align the textural style with its own internal data distribution. We denote this variant as LLaDA-rewrite-Qwen3.

3.3 Model Architecture and Training Strategy

Let (a_0, p_0, r_0) denote the audio input, textual prompt, and target response, respectively. We employ a frozen Whisper-small encoder to extract frame-level acoustic features from a_0 , and integrate them into the LLaDA-8B-Instruct backbone via two lightweight adapters.:

Semantic Adapter. A 2-layer convolutional network with a subsampling rate of 4, followed by a 2-layer linear projection. It compresses the encoder’s 50 Hz output to 12.5 Hz.

Acoustic Adapter. A 2-layer Q-former (Li et al., 2023) blocks with 64 trainable query vectors. It extracts speech-specific features from intermediate encoder states.

Two-Stage Training. The whole training process is shown in Figure 2. In stage 1, the semantic adapter is trained on 960 hours Librispeech using an ASR-style objective to align the speech encoder with the language model. In stage 2, both adapters are fine-tuned on our 127-hour synthetic dataset under the audio captioning objective. The final audio representation is the concatenation of outputs from both adapters and prepended as prefix tokens to the LLM input. In both stages, audio and prompt tokens remain unmasked during training. We use `<endofxt>` as a padding and end-of-sequence token during training, which must also be predicted. The LLaDA model and Whisper encoder remain frozen throughout. At each training step, the tokens of r_0 are independently replaced with a special mask token M with probability $t \in (0, 1]$. And then a forward masking process to sample a corrupted sequence r_t . We optimize the model using a diffusion-style masked prediction objective:




Figure 3: Inference procedure of DIFFA.

$$L_a = -\mathbb{E}_{t,a_0,p_0,r_0,r_t} \left[\frac{1}{t} \sum_{i=1}^{L'} \mathbf{1}[r_t^i = M] \log p_\theta(r_0^i | a_0, p_0, r_t) \right], \quad (3)$$

where r_t is the masked response and L' is its length.

3.4 Inference Procedure

At inference time, we first pad the prompt and audio input, then initialize the response r_T as a fully masked sequence of desired length. The model iteratively refines r_t over T denoising steps.

At step $t \rightarrow s$, the model predicts masked tokens:

$$\hat{r}_t = \arg \max p_\theta(r_0 | a_0, p_0, r_t), \quad (4)$$

then re-masks $\lceil s/t \rceil$ proportion of tokens with the lowest confidence to form r_s .

We follow a semi-autoregressive strategy (Nie et al., 2025), generating the sequence block-wise from left to right. Within each block, tokens are predicted in parallel and partially remasked.

This iterative inference scheme balances generation quality and efficiency, while maintaining the benefits of parallel decoding and bidirectional context modeling inherent in diffusion-based LLMs.

4 Experimental Setup

4.1 Datasets

In our experiments, we employ only open-source datasets Librispeech for ASR task in stage 1 and

Dataset	Samples	Total Dur. (h)
VCTK-Corpus	20,000	19.91
Accentdb	16,874	19.28
IEMOCAP	20,000	24.82
dailytalk	20,000	18.17
VoxCeleb1	20,000	45.83
Total	96,874	127.01

Table 1: Statistics of Datasets

five dataset to construct dataset for Q&A in stage 2: VCTK-Corpus (Yamagishi et al., 2019), Accentdb (Ahamad et al., 2020), IEMOCAP (Busso et al., 2008), dailytalk (Lee et al., 2023), Vox-Celeb (Nagrani et al., 2017). The details of datasets are presented in Table 1.

Compared to DESTA-2, our dataset follows a similar construction paradigm but excludes the *PromptTTS* and *Mixed Noise & Reverb* subsets due to lack of access. Our dataset includes 10 annotated attributes—*gender, age, accent, emotion, pitch, volume, speaking speed, duration, intent, and spoken text*—which is slightly fewer than the 12 attributes used in DESTA-2.

4.2 Model Configuration and Training Setup

For the speech encoder, we adopt the Whisper-Small encoder, which contains 88.2 million parameters. As the language backbone, we use LLaDA-8B-Instruct, a large language model trained with a masked denoising objective inspired by diffusion-based frameworks. It is built upon a Transformer decoder architecture with 32 layers, 32 attention heads, a hidden size of 4096, and approximately 8.1 billion parameters. The architecture follows LLaMA (Touvron et al., 2023; Dubey et al., 2024), with key modifications including RMSNorm (Zhang and Sennrich, 2019) for normalization, SwiGLU (Shazeer, 2020) for non-linearity, and rotary position embeddings (RoPE) (Su et al., 2024) for positional encoding.

In our experiments, all parameters of LLaDA-8B-Instruct are frozen. We introduce lightweight trainable adapters to integrate audio features, following prior work on parameter-efficient multimodal learning. The semantic adapter contains 14.4 million parameters and the acoustic adapter 22.3 million. Training details are provided in Appendix A.1.

Models	Perception				Reasoning				Overall Avg
	Semantics	Phonology	Paralinguistics	Avg	Semantics	Phonology	Paralinguistics	Avg	
Human	87.10	94.32	92.88	91.24	82.16	87.60	89.12	86.77	89.72
Gemini-1.5-Pro (Team et al., 2024)	57.06	53.60	31.23	46.10	79.47	83.46	46.33	76.16	60.68
Qwen2.5-Omni (Xu et al., 2025)	55.12	37.33	39.35	42.50	88.00	81.37	48.36	79.83	60.57
Kimi-Audio (Ding et al., 2025)	57.64	42.30	35.74	43.52	81.77	76.65	55.22	76.03	59.28
MiniCPM (Team, 2025)	56.56	34.05	36.48	40.54	80.71	74.72	46.71	73.57	56.53
GPT-4o-Audio (OpenAI et al., 2024)	59.70	41.56	21.44	39.67	80.83	78.74	26.25	71.96	56.38
MERaLiON (He et al., 2024)	54.49	33.69	25.84	35.74	80.32	77.18	41.49	73.68	54.10
Qwen2-Audio-Instruct (Chu et al., 2024)	52.14	32.87	35.56	39.02	77.62	64.81	46.67	68.90	53.27
Gemini-2.0-Flash	47.17	41.30	30.62	40.83	70.69	70.69	36.16	47.83	51.03
Megrez-3B-Omni (Li et al., 2025)	41.36	32.52	26.35	32.48	73.53	66.11	40.42	67.05	49.03
DIVA(Held et al., 2024)	44.36	33.72	27.45	33.95	62.32	74.24	40.00	65.04	48.31
Qwen-Audio-Chat (Chu et al., 2023)	57.21	38.52	24.70	35.69	58.61	59.78	25.60	55.93	46.92
Step-Audio (Huang et al., 2025)	31.56	29.39	24.01	28.72	49.10	50.09	45.27	47.27	37.42
BLSP (Wang et al., 2023)	31.35	20.96	23.75	28.36	47.91	42.31	42.08	44.97	35.96
GLM-4-Voice (Zeng et al., 2024)	27.80	24.52	27.34	26.18	46.10	48.16	44.35	46.76	35.51
Random	24.30	25.70	26.10	24.90	23.80	25.40	25.40	25.02	25.37
DIFFA	52.67	36.65	35.12	40.28	81.53	72.68	45.67	72.92	56.04

Table 2: Performance breakdown on the MMSU benchmark across perception and reasoning dimensions.

4.3 Benchmarks

MMSU (Wang et al., 2025) is a large-scale benchmark aimed at evaluating the perception and reasoning capabilities of SpeechLLMs in authentic spoken language scenarios. It consists of 5,000 carefully curated audio-question-answer triplets across 47 diverse tasks, covering a wide spectrum of linguistic and paralinguistic phenomena—including phonetics, prosody, semantics, emotion, and speaker traits. Tasks are categorized into two main dimensions: perception (e.g., intonation detection, disfluency recognition) and reasoning (e.g., sarcasm detection, code-switch QA), thereby enabling a comprehensive assessment of models’ ability to extract, interpret, and reason over fine-grained acoustic and linguistic cues.

MMAU (Sakshi et al., 2025) is a benchmark designed to evaluate advanced audio understanding through human-annotated multiple-choice questions paired with audio clips. It covers three core domains—speech, music, and environmental sounds—and targets 27 distinct skills that require complex reasoning and expert-level knowledge. Unlike benchmarks focused on low-level perception, MMAU emphasizes high-level cognitive abilities, challenging models to perform multi-step reasoning and knowledge retrieval grounded in audio inputs. In our experiments, we use the Test-mini split of MMAU for evaluation.

VoiceBench (Chen et al., 2024b) is a comprehensive benchmark designed to evaluate the capabilities of LLM-based voice assistants. It primarily consists of audio queries synthesized via text-to-speech (TTS) from existing text-based benchmarks, simulating realistic user interactions in spo-

ken form. It focuses on semantic understanding and supports structured evaluation along three dimensions: general knowledge, instruction following, and safety. These dimensions respectively assess the model’s ability to answer questions, follow spoken constraints, and reject harmful prompts. By converting textual tasks into audio, VoiceBench offers a targeted and rigorous testbed for audio-language models.

5 Experiments

5.1 Evaluation on MMSU

To assess the advanced reasoning abilities of our diffusion-based model, we evaluate DIFFA on the **MMSU** benchmark, which assesses fine-grained spoken language understanding across perception (semantics, phonology, paralinguistics) and reasoning dimensions. As shown in Table 2, our model achieves an average accuracy of 56.04%, highlighting its ability to handle a wide range of linguistically grounded tasks.

Although the overall performance trails top proprietary models like Gemini-1.5-Pro (60.68%), DIFFA outperforms many strong autoregressive baselines, such as Qwen2-Audio-Instruct (53.27%) and Gemini-2.0-Flash (51.03%). This competitive result, despite relying solely on synthetic supervision and lightweight adapters, supports the viability of diffusion-based approaches for nuanced speech understanding.

DIFFA performs particularly well on semantic reasoning tasks (81.53%), benefiting from its strong language modeling backbone and ASR-aligned speech encoder. However, like most mod-

Model	Sound	Music	Speech	Average
Gemini 2.5 Pro (Comanici et al., 2025)	75.08	68.26	71.47	71.60
Qwen2.5-Omni (Xu et al., 2025)	78.10	65.90	70.60	71.53
Phi-4-multimodal (Abouelenin et al., 2025)	65.47	64.37	67.27	65.70
Audio Flamingo 2 Reasoning (Ghosh et al., 2025)	75.98	74.25	43.54	64.59
GPT-4o Audio (OpenAI et al., 2024)	64.56	56.29	66.67	62.51
Audio Flamingo 2 (Ghosh et al., 2025)	71.47	70.96	44.74	62.39
Qwen2-Audio-Instruct (Chu et al., 2024)	67.27	56.29	55.26	59.61
GPT-4o mini Audio (OpenAI et al., 2024)	50.75	39.22	69.07	53.01
Gemini Pro v1.5 (Team et al., 2024)	56.75	49.40	58.55	52.97
M2UGen (Liu et al., 2023)	43.24	37.13	33.33	37.90
MusiLingo (Deng et al., 2024)	43.24	40.12	31.23	38.20
SALMONN (Tang et al., 2024b)	41.14	37.13	26.43	34.90
MuLLaMa (Liu et al., 2024)	33.03	32.34	17.42	27.60
GAMA-IT (Ghosh et al., 2024)	30.93	26.74	10.81	22.83
GAMA (Ghosh et al., 2024)	31.83	17.71	12.91	20.82
LTU (Gong et al., 2024)	20.42	15.97	15.92	17.44
Audio Flamingo Chat (Kong et al., 2024b)	25.23	17.66	6.91	16.60
DIFFA	46.25	43.41	59.46	49.71

Table 3: Evaluation results on the MMAU benchmark. Each model is assessed across three core audio domains: sound, music, and speech.

Model	AlpacaEval	CommonEval	SD-QA	MMSU*	OBQA	IFEval	AdvBench	Overall
GPT-4o-Audio (OpenAI et al., 2024)	4.78	4.49	75.50	80.25	89.23	76.02	98.65	86.43
Kimi-Audio (Ding et al., 2025)	4.46	3.97	63.12	62.17	83.52	61.10	100.00	76.93
Baichuan-Omni-1.5 (Li et al., 2024)	4.50	4.05	43.40	57.25	74.51	54.54	97.31	71.14
GLM-4-Voice (Zeng et al., 2024)	3.97	3.42	36.98	39.75	53.41	25.92	88.08	55.99
DiVA (Held et al., 2024)	3.67	3.54	57.06	25.76	25.49	39.16	98.27	55.70
Qwen2-Audio (Chu et al., 2024)	3.74	3.43	35.72	35.72	49.45	26.33	96.73	55.34
Step-Audio (Huang et al., 2025)	4.13	3.09	44.21	28.33	33.85	27.96	69.62	49.77
LLaMA-Omni (Fang et al., 2025)	3.70	3.46	39.69	25.93	27.47	14.87	11.35	37.50
VITA (Fu et al., 2024)	3.38	2.15	27.94	25.70	29.01	22.82	26.73	34.68
Slam-Omni (Chen et al., 2024a)	1.90	1.79	4.16	26.06	25.27	13.38	94.23	33.84
Mini-Omni2 (Xie and Wu, 2024b)	2.32	2.18	9.31	24.27	26.59	11.56	57.50	31.32
Mini-Omni (Xie and Wu, 2024a)	1.95	2.02	13.92	24.69	26.59	13.58	37.12	27.90
Moshi (Défossez et al., 2024b)	2.01	1.60	15.64	24.04	25.93	10.12	44.23	27.45
DIFFA	3.78	2.96	34.45	29.57	35.60	26.56	76.54	48.22

Table 4: Evaluation results on VoiceBench. Metrics cover diverse QA and alignment tasks. Note that MMSU* in VoiceBench is derived from MMLU-Pro, which differs from the MMSU benchmark.

els in the benchmark, it exhibits lower accuracy in phonological and paralinguistic tasks—areas that demand precise acoustic perception beyond textual semantics. These trends mirror the broader challenges outlined in the MMSU benchmark, where human-level performance (89.72%) remains a distant target.

This result provides the first empirical evidence that diffusion-based language models can serve as viable backbones for large-scale audio-language understanding, even without autoregressive decoding. Despite using only 127 hours of synthetic training data—orders of magnitude less than the tens of thousands of supervised fine-tuning (SFT) hours used by many baselines—DIFFA demon-

strates strong generalization and reasoning capabilities.

Overall, DIFFA demonstrates robust generalization across complex linguistic dimensions, establishing a strong diffusion-based baseline. Future work should focus on enhancing the model’s sensitivity to prosodic and phonological cues to bridge the gap with human-level performance.

5.2 Evaluation on MMAU

We further evaluate DIFFA on the MMAU benchmark, which tests 27 skills across three audio domains: sound, music, and speech. As shown in Table 3, DIFFA achieves an average accuracy of 49.71%, outperforming several widely used autore-

LLM Backbone	Data Source	Adapter	MMAU				MMSU		
			Sound	Music	Speech	Avg	Perception	Reasoning	Avg
LLaMA 3.1	LLaMA 3	Dual	22.82	26.65	35.74	28.40	32.36	44.83	38.40
LLaDA	LLaMA 3	Dual	47.75	45.51	61.86	51.71	37.27	73.20	54.72
LLaDA	Qwen3	Single	44.44	44.31	54.35	47.70	36.98	69.83	52.88
LLaDA	Qwen3	Dual	46.25	43.41	59.46	49.71	40.28	72.92	56.04

Table 5: Ablation study on model architecture and adapter design. All models use 8B Instruct version.

Data Source	MMAU	MMSU	Voicebench	Avg
LLaMA 3	51.71	54.72	37.17	47.86
LLaDA	51.31	56.18	43.52	50.34
Qwen3	49.71	56.04	48.22	51.32
rewrite-Qwen3	50.41	56.43	46.60	51.15

Table 6: Ablation study on the impact of different instruction data sources. All models use 8B Instruct version.

gressive LALMs, such as SALMONN (34.90%), GAMA-IT (22.83%), and LTU (17.44%). It also approaches the performance of commercial models like GPT-4o mini Audio (53.01%) and Gemini Pro v1.5 (52.97%).

A domain-level breakdown shows that DIFFA achieves the highest performance on speech-related tasks (59.46%), likely benefiting from its speech-caption-focused training paradigm. In contrast, models such as Audio Flamingo 2 perform better on music and environmental sounds but underperform in speech understanding (e.g., 44.74% for speech), underscoring the advantage of targeted, speech-centric training.

Overall, these results position DIFFA as a promising diffusion-based alternative to autoregressive LALMs. With limited resources and a parameter-efficient adapter tuning scheme, it achieves competitive results across complex audio reasoning benchmarks, indicating strong potential for scaling with larger or higher-quality datasets.

5.3 Evaluation on VoiceBench

VoiceBench evaluates semantic understanding from audio-as-question prompts, covering knowledge, instructions, and safety—making it a rigorous benchmark for spoken query comprehension.

Despite being trained on only 960 hours of ASR data and 127 hours of synthetic instructions, DIFFA achieves 34.45% on SD-QA and 35.60% on OBQA, demonstrating promising capability in factual spoken QA. This is particularly notable when compared to models like Qwen2-Audio (35.72% SD-

QA, 49.45% OBQA), which are trained on hundreds of thousands of hours of proprietary data. The results suggest that diffusion-based models, even with limited training, can capture core semantic structures in speech.

On IFEval, DIFFA reaches 26.56%, slightly surpassing Qwen2-Audio (26.33%) and GLM-4-Voice (25.92%), indicating a basic capacity for instruction comprehension from audio inputs. However, the performance gap with top models such as Kimi-Audio (61.10%) underscores the challenge of aligning audio-conditioned instruction execution without large-scale supervised tuning.

In AdvBench, DIFFA attains 76.54%, outperforming many strong baselines and approaching GLM-4-Voice (88.08). This highlights the potential of lightweight, diffusion-based models to learn safety-aligned behavior with minimal data.

In summary, DIFFA establishes a competitive baseline on VoiceBench despite using orders of magnitude less training data than competing systems. These findings validate the feasibility of diffusion-based LLMs for semantic audio understanding and offer a data-efficient alternative to current autoregressive paradigms.

5.4 Ablation Study

We perform a comprehensive ablation study to assess the impact of three key factors on model performance: (1) the choice of language model backbone, (2) adapter design, and (3) instruction data source. Besides, effect of inference hyperparameters are provided in Appendix A.2.

Impact of Diffusion-Based Language Modeling.

As shown in Table 5, replacing the autoregressive LLaMA 3.1 backbone with the diffusion-based LLaDA architecture leads to a substantial improvement across all metrics. Specifically, with dual adapters, the LLaDA variant achieves 51.71 on MMAU and 54.72 on MMSU, significantly outperforming its LLaMA counterpart. This highlights the advantage of diffusion language model for au-

audio understanding tasks.

Effect of Adapter Design. We compare single and dual adapter configurations using the same LLaDA backbone. The single adapter setup uses only a semantic adapter aligned with the speech encoder’s output. The dual adapter adds an acoustic adapter that extracts low-level acoustic cues from intermediate encoder states. Results show that dual adapters yield a +2.01 gain on MMAU and +3.16 on MMSU, confirming the advantage of combining semantic and acoustic information for richer audio understanding.

Instruction Data Source. Table 6 examines the role of different instruction data sources on final performance. All variants show comparable results, with Qwen3-generated instruction data leading to the best overall accuracy. Rewriting Qwen3 data with LLaDA brings only marginal improvements, suggesting that enhancing data quality at generation time may be more effective than post-hoc refinement. Interestingly, models trained on LLaDA-generated data outperform those based on LLaMA-3 instructions, indicating that the inductive bias of diffusion-based generation may yield more aligned supervision.

6 Conclusion

In this work, we introduce **DIFFA**, a diffusion-based large audio-language model that combines a frozen diffusion language backbone with lightweight dual adapters for audio understanding and instruction following. Using a two-stage training strategy and fully synthetic instruction data, DIFFA achieves competitive performance on diverse benchmarks—including MMSU, MMAU, and VoiceBench—despite relying on only 960 hours of ASR data and 127 hours of synthetic instruction data, in contrast to models like Qwen2-Audio-Instruct trained on over 500K hours. As an initial exploration of diffusion-based modeling in the audio domain, our results suggest that such models offer a promising alternative to autoregressive LALMs. While our experiments are conducted on relatively small-scale open-source corpora, we plan to scale to broader and more diverse data in future work. We hope this study encourages further research into efficient, flexible, and controllable speech-driven AI systems.

Limitations

A notable limitation is the constrained volume of training data. DIFFA is trained on a modest dataset comprising 960 hours of ASR data and 127 hours of synthetic instruction data, which may limit its capacity to generalize to a broader range of real-world audio phenomena—such as low-resource accents, noisy acoustic environments, or specialized domain speech. To address this, future iterations will focus on scaling up the training data, including expanding both ASR corpora and diverse synthetic instruction datasets, to strengthen the model’s robustness and coverage of complex audio scenarios.

References

- Abdelrahman Abouelenin, Atabak Ashfaq, Adam Atkinson, Hany Awadalla, Nguyen Bach, Jianmin Bao, Alon Benhaim, Martin Cai, Vishrav Chaudhary, Congcong Chen, et al. 2025. Phi-4-mini technical report: Compact yet powerful multimodal language models via mixture-of-loras. *arXiv preprint arXiv:2503.01743*.
- Josh Achiam, Steven Adler, Sandhini Agarwal, Lama Ahmad, Ilge Akkaya, Florencia Leoni Aleman, Diogo Almeida, Janko Altschmidt, Sam Altman, Shyamal Anadkat, et al. 2023. Gpt-4 technical report. *arXiv preprint arXiv:2303.08774*.
- Afroz Ahamad, Ankit Anand, and Pranesh Bhargava. 2020. **AccentDB: A database of non-native English accents to assist neural speech recognition**. In *Proceedings of the Twelfth Language Resources and Evaluation Conference*, pages 5351–5358, Marseille, France. European Language Resources Association.
- Jacob Austin, Daniel D. Johnson, Jonathan Ho, Daniel Tarlow, and Rianne Vandenberg. 2021. Structured denoising diffusion models in discrete state-spaces. *arXiv: Learning, arXiv: Learning*.
- Carlos Busso, Murtaza Bulut, Chi-Chun Lee, Abe Kazemzadeh, Emily Mower, Samuel Kim, Jeanette N. Chang, Sungbok Lee, and Shrikanth S. Narayanan. 2008. **Iemocap: interactive emotional dyadic motion capture database**. *Language Resources and Evaluation*, page 335–359.
- Wenxi Chen, Ziyang Ma, Ruiqi Yan, Yuzhe Liang, Xiquan Li, Ruiyang Xu, Zhikang Niu, Yanqiao Zhu, Yifan Yang, Zhanxun Liu, et al. 2024a. Slam-omni: Timbre-controllable voice interaction system with single-stage training. *arXiv preprint arXiv:2412.15649*.
- Yiming Chen, Xianghu Yue, Chen Zhang, Xiaoxue Gao, Robby T Tan, and Haizhou Li. 2024b. Voicebench: Benchmarking llm-based voice assistants. *arXiv preprint arXiv:2410.17196*.

- Yunfei Chu, Jin Xu, Qian Yang, Haojie Wei, Xipin Wei, Zhifang Guo, Yichong Leng, Yuanjun Lv, Jinzheng He, Junyang Lin, Chang Zhou, and Jingren Zhou. 2024. Qwen2-audio technical report. *arXiv preprint arXiv:2407.10759*.
- Yunfei Chu, Jin Xu, Xiaohuan Zhou, Qian Yang, Shiliang Zhang, Zhijie Yan, Chang Zhou, and Jingren Zhou. 2023. Qwen-audio: Advancing universal audio understanding via unified large-scale audio-language models. *arXiv preprint arXiv:2311.07919*.
- Gheorghe Comanici, Eric Bieber, Mike Schaekermann, Ice Pasupat, Naveen Sachdeva, Inderjit Dhillon, Marcel Blistein, Ori Ram, Dan Zhang, Evan Rosen, et al. 2025. Gemini 2.5: Pushing the frontier with advanced reasoning, multimodality, long context, and next generation agentic capabilities. *arXiv preprint arXiv:2507.06261*.
- Alexandre Défossez, Laurent Mazaré, Manu Orsini, Amélie Royer, Patrick Pérez, Hervé Jégou, Edouard Grave, and Neil Zeghidour. 2024a. Moshi: a speech-text foundation model for real-time dialogue. *arXiv preprint arXiv:2410.00037*.
- Alexandre Défossez, Laurent Mazaré, Manu Orsini, Amélie Royer, Patrick Pérez, Hervé Jégou, Edouard Grave, and Neil Zeghidour. 2024b. Moshi: a speech-text foundation model for real-time dialogue. *arXiv preprint arXiv:2410.00037*.
- Zihao Deng, Yinghao Ma, Yudong Liu, Rongchen Guo, Ge Zhang, Wenhui Chen, Wenhao Huang, and Emmanuel Benetos. 2024. [Musilingo: Bridging music and text with pre-trained language models for music captioning and query response](#). In *NAACL-HLT (Findings)*, pages 3643–3655.
- Ding Ding, Zeqian Ju, Yichong Leng, Songxiang Liu, Tong Liu, Zeyu Shang, Kai Shen, Wei Song, Xu Tan, Heyi Tang, et al. 2025. Kimi-audio technical report. *arXiv preprint arXiv:2504.18425*.
- Abhimanyu Dubey, Abhinav Jauhri, Abhinav Pandey, Abhishek Kadian, Ahmad Al-Dahle, Aiesha Letman, Akhil Mathur, Alan Schelten, Amy Yang, Angela Fan, et al. 2024. The llama 3 herd of models. *arXiv e-prints*, pages arXiv–2407.
- Qingkai Fang, Shoutao Guo, Yan Zhou, Zhengrui Ma, Shaolei Zhang, and Yang Feng. 2025. [LLaMA-omni: Seamless speech interaction with large language models](#). In *The Thirteenth International Conference on Learning Representations*.
- Yassir Fathullah, Chunyang Wu, Egor Lakomkin, Ke Li, Junteng Jia, Yuan Shangquan, Jay Mahadeokar, Ozlem Kalinli, Christian Fuegen, and Mike Seltzer. 2024. [AudioChatLlama: Towards general-purpose speech abilities for LLMs](#). In *Proceedings of the 2024 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 1: Long Papers)*, pages 5522–5532, Mexico City, Mexico. Association for Computational Linguistics.
- Chaoyou Fu, Haojia Lin, Zuwei Long, Yunhang Shen, Yuhang Dai, Meng Zhao, Yi-Fan Zhang, Shaoqi Dong, Yangze Li, Xiong Wang, et al. 2024. Vita: Towards open-source interactive omni multimodal llm. *arXiv preprint arXiv:2408.05211*.
- Sreyan Ghosh, Zhifeng Kong, Sonal Kumar, S Sakshi, Jaehyeon Kim, Wei Ping, Rafael Valle, Dinesh Manocha, and Bryan Catanzaro. 2025. [Audio flamingo 2: An audio-language model with long-audio understanding and expert reasoning abilities](#). In *Forty-second International Conference on Machine Learning*.
- Sreyan Ghosh, Sonal Kumar, Ashish Seth, Chandra Kiran Reddy Evuru, Utkarsh Tyagi, S. Sakshi, Oriol Nieto, Ramani Duraiswami, and Dinesh Manocha. 2024. [Gama: A large audio-language model with advanced audio understanding and complex reasoning abilities](#). In *EMNLP*, pages 6288–6313.
- Yuan Gong, Hongyin Luo, Alexander H. Liu, Leonid Karlinsky, and James R. Glass. 2024. [Listen, think, and understand](#). In *The Twelfth International Conference on Learning Representations*.
- Yingxu He, Zhuohan Liu, Shuo Sun, Bin Wang, Wenyu Zhang, Xunlong Zou, Nancy F Chen, and Ai Ti Aw. 2024. Meralion-audiollm: Bridging audio and language with large language models. *arXiv preprint arXiv:2412.09818*.
- William Held, Ella Li, Michael Ryan, Weiyan Shi, Yanzhe Zhang, and Diyi Yang. 2024. Distilling an end-to-end voice assistant without instruction training data. *arXiv preprint arXiv:2410.02678*.
- Ailin Huang, Boyong Wu, Bruce Wang, Chao Yan, Chen Hu, Chengli Feng, Fei Tian, Feiyu Shen, Jingbei Li, Mingrui Chen, et al. 2025. Step-audio: Unified understanding and generation in intelligent speech interaction. *arXiv preprint arXiv:2502.11946*.
- Zhifeng Kong, Arushi Goel, Rohan Badlani, Wei Ping, Rafael Valle, and Bryan Catanzaro. 2024a. Audio flamingo: A novel audio language model with few-shot learning and dialogue abilities. *arXiv preprint arXiv:2402.01831*.
- Zhifeng Kong, Arushi Goel, Rohan Badlani, Wei Ping, Rafael Valle, and Bryan Catanzaro. 2024b. Audio flamingo: A novel audio language model with few-shot learning and dialogue abilities. In *International Conference on Machine Learning*, pages 25125–25148. PMLR.
- Keon Lee, Kyumin Park, and Daeyoung Kim. 2023. [Dailytalk: Spoken dialogue dataset for conversational text-to-speech](#). In *ICASSP 2023 - 2023 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*, pages 1–5.
- Boxun Li, Yadong Li, Zhiyuan Li, Congyi Liu, Weilin Liu, Guowei Niu, Zheyue Tan, Haiyang Xu, Zhuyu Yao, Tao Yuan, et al. 2025. Megrez-omni technical report. *arXiv preprint arXiv:2502.15803*.

- Junnan Li, Dongxu Li, Silvio Savarese, and Steven Hoi. 2023. Blip-2: Bootstrapping language-image pre-training with frozen image encoders and large language models. In *International conference on machine learning*, pages 19730–19742. PMLR.
- Yadong Li, Haoze Sun, Mingan Lin, Tianpeng Li, Guosheng Dong, Tao Zhang, Bowen Ding, Wei Song, Zhenglin Cheng, Yuqi Huo, et al. 2024. Baichuan-omni technical report. *arXiv preprint arXiv:2410.08565*.
- Shansong Liu, Atin Sakkeer Hussain, Chenshuo Sun, and Ying Shan. 2023. M2ugen: Multi-modal music understanding and generation with the power of large language models. *arXiv preprint arXiv:2311.11255*.
- Shansong Liu, Atin Sakkeer Hussain, Chenshuo Sun, and Ying Shan. 2024. Music understanding llama: Advancing text-to-music generation with question answering and captioning. In *ICASSP 2024-2024 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*, pages 286–290. IEEE.
- Ke-Han Lu, Zhehuai Chen, Szu-Wei Fu, Chao-Han Huck Yang, Jagadeesh Balam, Boris Ginsburg, Yu-Chiang Frank Wang, and Hung-yi Lee. 2025. Developing instruction-following speech language model without speech instruction-tuning data. In *ICASSP 2025-2025 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*, pages 1–5. IEEE.
- Arsha Nagrani, Joon Son Chung, and Andrew Senior. 2017. *Voxceleb: a large-scale speaker identification dataset*. In *Interspeech 2017*.
- Shen Nie, Fengqi Zhu, Zebin You, Xiaolu Zhang, Jingyang Ou, Jun Hu, Jun Zhou, Yankai Lin, Ji-Rong Wen, and Chongxuan Li. 2025. Large language diffusion models. *arXiv preprint arXiv:2502.09992*.
- OpenAI, :, Aaron Hurst, Adam Lerer, Adam P. Goucher, Adam Perelman, Aditya Ramesh, Aidan Clark, AJ Ostrow, Akila Welihinda, Alan Hayes, Alec Radford, Aleksander Madry, Alex Baker-Whitcomb, Alex Beutel, Alex Borzunov, Alex Carney, Alex Chow, Alex Kirillov, Alex Nichol, Alex Paino, Alex Renzin, Alex Tachard Passos, Alexander Kirillov, Alexi Christakis, Alexis Conneau, Ali Kamali, Allan Jabri, Allison Moyer, Allison Tam, Amadou Crookes, Amin Tootoochian, Amin Tootoonchian, Ananya Kumar, and et al. 2024. *Gpt-4o system card*. *Preprint*, arXiv:2410.21276.
- Jingyang Ou, Shen Nie, Kaiwen Xue, Fengqi Zhu, Jiacheng Sun, Zhenguo Li, and Chongxuan Li. 2025. *Your absorbing discrete diffusion secretly models the conditional distributions of clean data*. In *The Thirteenth International Conference on Learning Representations*.
- Vassil Panayotov, Guoguo Chen, Daniel Povey, and Sanjeev Khudanpur. 2015. Librispeech: an asr corpus based on public domain audio books. In *2015 IEEE international conference on acoustics, speech and signal processing (ICASSP)*, pages 5206–5210. IEEE.
- Alec Radford, Jong Wook Kim, Tao Xu, Greg Brockman, Christine McLeavey, and Ilya Sutskever. 2023. Robust speech recognition via large-scale weak supervision. In *International conference on machine learning*, pages 28492–28518. PMLR.
- Subham Sekhar Sahoo, Marianne Arriola, Yair Schiff, Aaron Gokaslan, Edgar Marroquin, Justin T Chiu, Alexander Rush, and Volodymyr Kuleshov. 2024. *Simple and effective masked diffusion language models*. *ArXiv*, abs/2406.07524.
- S Sakshi, Utkarsh Tyagi, Sonal Kumar, Ashish Seth, Ramaneswaran Selvakumar, Oriol Nieto, Ramani Duraiswami, Sreyan Ghosh, and Dinesh Manocha. 2025. *MMAU: A massive multi-task audio understanding and reasoning benchmark*. In *The Thirteenth International Conference on Learning Representations*.
- Noam Shazeer. 2020. Glu variants improve transformer. *arXiv preprint arXiv:2002.05202*.
- Jiaxin Shi, Kehang Han, Zhe Wang, Arnaud Doucet, and Michalis Titsias. 2024. Simplified and generalized masked diffusion for discrete data. *Advances in neural information processing systems*, 37:103131–103167.
- Jianlin Su, Murtadha Ahmed, Yu Lu, Shengfeng Pan, Wen Bo, and Yunfeng Liu. 2024. Roformer: Enhanced transformer with rotary position embedding. *Neurocomputing*, 568:127063.
- Changli Tang, Wenyi Yu, Guangzhi Sun, Xianzhao Chen, Tian Tan, Wei Li, Lu Lu, Zejun MA, and Chao Zhang. 2024a. *SALMONN: Towards generic hearing abilities for large language models*. In *The Twelfth International Conference on Learning Representations*.
- Changli Tang, Wenyi Yu, Guangzhi Sun, Xianzhao Chen, Tian Tan, Wei Li, Lu Lu, Zejun Ma, and Chao Zhang. 2024b. *Salmonn: Towards generic hearing abilities for large language models*. In *ICLR*.
- Gemini Team, Petko Georgiev, Ving Ian Lei, Ryan Burnell, Libin Bai, Anmol Gulati, Garrett Tanzer, Damien Vincent, Zhufeng Pan, Shibo Wang, et al. 2024. Gemini 1.5: Unlocking multimodal understanding across millions of tokens of context. *arXiv preprint arXiv:2403.05530*.
- OpenBMB MiniCPM-o Team. 2025. Minicpm-o 2.6: A gpt-4o level mllm for vision, speech, and multimodal live streaming on your phone.
- Hugo Touvron, Louis Martin, Kevin Stone, Peter Albert, Amjad Almahairi, Yasmine Babaei, Nikolay Bashlykov, Soumya Batra, Prajjwal Bhargava, Shrutu Bhosale, et al. 2023. Llama 2: Open foundation and fine-tuned chat models. *arXiv preprint arXiv:2307.09288*.

Chen Wang, Minpeng Liao, Zhongqiang Huang, Jinliang Lu, Junhong Wu, Yuchen Liu, Chengqing Zong, and Jiajun Zhang. 2023. Blsp: Bootstrapping language-speech pre-training via behavior alignment of continuation writing. *arXiv preprint arXiv:2309.00916*.

Dingdong Wang, Jincenzi Wu, Junan Li, Dongchao Yang, Xueyuan Chen, Tianhua Zhang, and Helen Meng. 2025. Mmsu: A massive multi-task spoken language understanding and reasoning benchmark. *arXiv preprint arXiv:2506.04779*.

Zhifei Xie and Changqiao Wu. 2024a. Mini-omni: Language models can hear, talk while thinking in streaming. *arXiv preprint arXiv:2408.16725*.

Zhifei Xie and Changqiao Wu. 2024b. Mini-omni2: Towards open-source gpt-4o with vision, speech and duplex capabilities. *arXiv preprint arXiv:2410.11190*.

Jin Xu, Zhifang Guo, Jinzheng He, Hangrui Hu, Ting He, Shuai Bai, Keqin Chen, Jialin Wang, Yang Fan, Kai Dang, et al. 2025. Qwen2. 5-omni technical report. *arXiv preprint arXiv:2503.20215*.

Junichi Yamagishi, Christophe Veaux, and Kirsten Macdonald. 2019. Cstr vctk corpus: English multi-speaker corpus for cstr voice cloning toolkit (version 0.92).

Zhaorui Yang, Tianyu Pang, Haozhe Feng, Han Wang, Wei Chen, Minfeng Zhu, and Qian Liu. 2024. [Self-distillation bridges distribution gap in language model fine-tuning](#). In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 1028–1043, Bangkok, Thailand. Association for Computational Linguistics.

Zebin You, Shen Nie, Xiaolu Zhang, Jun Hu, Jun Zhou, Zhiwu Lu, Ji-Rong Wen, and Chongxuan Li. 2025. Llada-v: Large language diffusion models with visual instruction tuning. *arXiv preprint arXiv:2505.16933*.

Aohan Zeng, Zhengxiao Du, Mingdao Liu, Kedong Wang, Shengmin Jiang, Lei Zhao, Yuxiao Dong, and Jie Tang. 2024. Glm-4-voice: Towards intelligent and human-like end-to-end spoken chatbot. *arXiv preprint arXiv:2412.02612*.

Biao Zhang and Rico Sennrich. 2019. Root mean square layer normalization. *Advances in neural information processing systems*, 32.

Dong Zhang, Shimin Li, Xin Zhang, Jun Zhan, Pengyu Wang, Yaqian Zhou, and Xipeng Qiu. 2023. Speechgpt: Empowering large language models with intrinsic cross-modal conversational abilities. *arXiv preprint arXiv:2305.11000*.

A Hyperparameter Settings

A.1 Training Hyperparameter

In Stage 1, we train the semantic adapter using the LibriSpeech dataset for 10 epochs. We adopt a

learning rate of $1e^{-4}$ with 1000 warm-up steps and a global batch size of 128. In Stage 2, both the semantic and acoustic adapters are jointly trained on our generated dataset for 10 epochs. We use a learning rate of $5e^{-5}$ with 2000 warm-up steps and a global batch size of 32. All experiments are optimized using the Adam optimizer and conducted on 4 NVIDIA A800 GPUs with 80GB memory each.

A.2 Inference Hyperparameter

Benchmark	Answer length	Block length	Steps
MMSU	4	4	4
MMAU	16	16	16
AlpacaEval	128	32	128
CommonEval	128	32	128
SD-QA	128	32	128
MMSU	16	16	16
OBQA	16	16	16
IFEval	256	32	256
AdvBench	128	32	128

Table A.1: Inference hyperparameters used for each benchmark

We detail the inference hyperparameters used in our experiments across different benchmarks in Table A.1. Specifically, we configure the answer length (i.e., maximum number of tokens generated), block length (i.e., the number of tokens decoded per block from left to right), and the total number of denoising steps used in the diffusion-based decoding. For further details on the decoding process, please refer to LLaDA (Nie et al., 2025). For decoding steps, we set it equal to answer length to achieve the best performance.

Answer length	Block length	Steps	AlpacaEval
128	16	128	3.72
128	32	128	3.78
128	64	128	3.56
128	128	128	3.07
256	16	256	3.78
256	32	256	3.76
256	64	256	3.65
256	128	256	3.28

Table A.2: Effect of answer length and block length on AlpacaEval performance

Effect of answer length and block length on AlpacaEval performance are reported in Table A.2. We observe that performance is sensitive to both the total answer length and the block size. Moderate block sizes (e.g., 32) consistently yield better

results, likely due to a balance between context stability and denoising granularity.